

Cache Performance Improvement

Dr A Sahu

A Sahu

1

Eleven Advanced Optimization for Cache Performance

- Reducing hit time
- Reducing miss penalty
- Reducing miss rate
- Reducing miss penalty * miss rate

Ref: 5.2, **Computer Architecture: A Quantitative Approach**, *Hennessy Patterson* Book, 4th Edition, PDF Version Available on Course website (Intranet)

A Sahu

2

Reducing Cache Hit Time

A Sahu

3

Reducing Hit Time

- Small and simple caches
- Pipelined cache access
- Trace caches
- Avoid time loss in address translation
 - Virtually indexed, physically tagged cache
 - simple and effective approach
 - possible only if cache is not too large
 - Virtually addressed cache
 - protection?, multiple processes?, aliasing?, I/O?

A Sahu

4


Small and Simple Caches

- Small size => faster access
- Small size => fit on the chip, lower delay
- Simple (direct mapped) => lower delay
- Second level – tags may be kept on chip

A Sahu

5

Cache access time estimates using CACTI



.8 micron technology, 1 R/W port, 32 b address, 64 b o/p, 32 B block

© 2003 Elsevier Science (USA). All rights reserved.

A Sahu

6

Pipelined Cache Access

- Multi-cycle cache access but pipelined
- reduces cycle time but hit time is more than one cycle
- Pentium 4 takes 4 cycles
- Greater penalty on branch misprediction
- More clock cycles between issue of load and use of data

— IF IF IF in pipeline

A Sahu

7

Trace Caches: Pre-decoded

- What maps to a cache block?
 - not statically determined
 - decided by the dynamic sequence of instructions, including predicted branches
- Used in Pentium 4 (NetBurst architecture)
- starting addresses not *word size * powers of 2*
- Better utilization of cache space
- downside – same instruction may be stored multiple times

A Sahu

8

Reducing Cache Miss Penalty

A Sahu

9

Reducing Miss Penalty

- Multi level caches
- Critical word first and early restart
- Giving priority to read misses over write
- Merging write buffer
- Victim caches

A Sahu

10

Multi Level Caches

Average memory access time =

$$\text{Hit time}_{L1} + \text{Miss rate}_{L1} * \text{Miss penalty}_{L1}$$

Miss penalty_{L1} =

$$\text{Hit time}_{L2} + \text{Miss rate}_{L2} * \text{Miss penalty}_{L2}$$

A Sahu

11

Critical Word First and Early Restart

- Read policy: Concurrent Read/Forward
- Load policy: Wrap around load

More effective when block size is large

A Sahu

12

Read Miss Priority Over Write

- Provide write buffers
- Processor writes into buffer and proceeds (for write through as well as write back)

On read miss

- wait for buffer to be empty, or
- check addresses in buffer for conflict

A Sahu

13

Merging Write Buffer


Merge writes belonging to same block in case of write through

A Sahu

14

Victim Cache: Recycle bin/Dust bin

- Evicted blocks are recycled
- Much faster than getting a block from the next level
- Size = 1 to 5 blocks
- A significant fraction of misses may be found in victim cache



A Sahu

15

Reducing Cache Miss Rate

A Sahu

16

Reducing Miss Rate

- Large Block Size
- Larger Cache
- Higher Associativity
- Way prediction and pseudo-associative cache
- Compiler optimizations

A Sahu

17

Large Block Size

- Take benefit of spatial locality
- Reduces compulsory misses
- Too large block size - misses increase
- Miss Penalty increases

A Sahu

18

Large Cache

- Reduces capacity misses
- **Hit time increases**
- Keep small L1 cache and large L2 cache

A Sahu

19

Higher Associativity

- Reduces conflict misses
- **8-way is almost like fully associative**
- **Hit time increases: What to do ?**
– **Pseudo Associativity**

A Sahu

20

Way Prediction and Pseudo-associative Cache

Way prediction: low miss rate of SA cache with hit time of DM cache

- Only one tag is compared initially
- Extra bits are kept for prediction
- Hit time in case of mis-prediction is high

Pseudo-assoc. or column assoc. cache: get advantage of SA cache in a DM cache

- Check sequentially in a pseudo-set
- Fast hit and slow hit

A Sahu

21

Compiler optimizations

- **Loop interchange**
– Improve spatial locality by scanning arrays row-wise
- **Blocking**
– Improve temporal and spatial locality

A Sahu

22

Improving Locality

Matrix Multiplication example

$$[C] = [A] \times [B]$$

$L \times M$ $L \times N$ $N \times M$

A Sahu

23

Cache Organization for the example

- Cache line (or block) = 4 matrix elements.
- Matrices are stored row wise.
- **Cache can't accommodate a full row/column.**
– **L, M and N are so large w.r.t. the cache size**
– After an iteration along any of the three indices, when an element is accessed again, it results in a miss.
- Ignore misses due to conflict between matrices.
– As if there was a **separate cache for each matrix**.

A Sahu

24

Matrix Multiplication : Code I

```

for (i = 0; i < L; i++)
  for (j = 0; j < M; j++)
    for (k = 0; k < N; k++)
      C[i][j] += A[i][k] * B[k][j];

```

	C	A	B
accesses	LM	LMN	LMN
misses	LM/4	LMN/4	LMN

Total misses = $LM(5N+1)/4$

A Sahu

25

Matrix Multiplication : Code II

```

for (k = 0; k < N; k++)
  for (i = 0; i < L; i++)
    for (j = 0; j < M; j++)
      C[i][j] += A[i][k] * B[k][j];

```

	C	A	B
accesses	LMN	LN	LMN
misses	LMN/4	LN	LMN/4

Total misses = $LN(2M+4)/4$

A Sahu

26

Matrix Multiplication : Code III

```

for (i = 0; i < L; i++)
  for (k = 0; k < N; k++)
    for (j = 0; j < M; j++)
      C[i][j] += A[i][k] * B[k][j];

```

	C	A	B
accesses	LMN	LN	LMN
misses	LMN/4	LN/4	LMN/4

Total misses = $LN(2M+1)/4$

A Sahu

27

Reducing MissRate*MissPenalty

A Sahu

28

Reducing Miss Penalty * Miss Rate

- Non-blocking cache
- Hardware prefetching
- Compiler controlled prefetching

A Sahu

29

Non-blocking Cache**In OOO processor**

- Hit under a miss
 - complexity of cache controller increases
- Hit under multiple misses or miss under a miss
 - memory should be able to handle multiple misses

A Sahu

30


Hardware Prefetching

- **Prefetch items before they are requested**
 - both data and instructions
- **What and when to prefetch?**
 - fetch two blocks on a miss (requested+next)
- **Where to keep prefetched information?**
 - in cache
 - in a separate buffer (most common case)

A Sahu

31

Prefetch Buffer/Stream Buffer



A Sahu

32

MatMul: Code III

```
for (i = 0; i < L; i++)
  for (k = 0; k < N; k++)
    for (j = 0; j < M; j++)
      C[i][j] += A[i][k] * B[k][j];
```

	C	A	B
accesses	LMN	LN	LMN
misses	LMN/4	LN/4	LMN/4
Total misses = LN(2M+1)/4			

Suppose 3 Separate Prefetcher for A, B and C
 All the 3 block can be brought to buffer & one
 swap out in T_e , $T_e = 4$ time execution of stmt;
 Over
 How many number Miss ?

$$3=1+1+1$$

A Sahu

33

Compiler Controlled Pre-fetching

- Semantically invisible (no change in registers or cache contents)
- Makes sense if processor doesn't stall while prefetching (non-blocking cache)
- Overhead of prefetch instruction should not exceed the benefit

```
PreFecth(A[i]); // Prefetch
STMT;
STMT;
STMT;
Y=A[i]; // Using Data
Z=Y+K;
```

```
X= A[i]; // Prefetch Instr
STMT;
STMT;
STMT;
Y=A[i]; // Using Data
Z=Y+K;
```

34

Course after Mid: Multicore Computing



**Good luck for your
Mid Semester Exam**

A Sahu

35